



Introduction to Neural Nets

CGnal S.r.l – Corso Venezia 43 - Milano

19 Novembre 2021 | Milano

DAY 1

Introduction

- Brief overview of Machine Learning (Supervised, Unsupervised)
- Introduction to Graph, Graph Theory and main metrics for characterizing graphs

DAY 2

Graph Machine Learning

- Community detection on Graphs
- Supervised Machine Learning on Graphs

DAY 3

Explainability & Interpretability

- Introduction to explainability problem
- LIME & SHAP

DAY 4

Simple Neural Networks


- Introduction to Neural Networks, TensorFlow and Computational Graphs
- Implementation and training of simple Neural Networks

DAY 5

Advanced Neural Networks

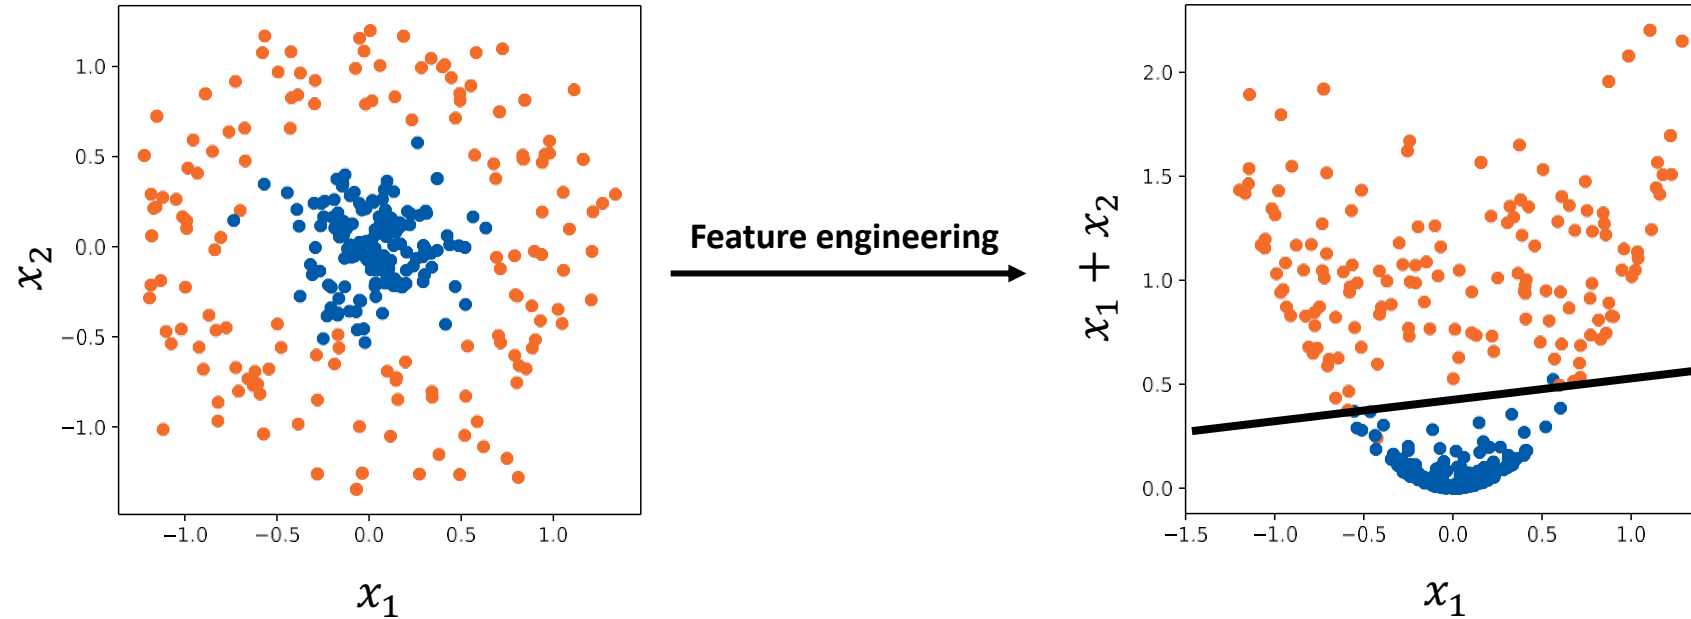
- Convolutional Neural Networks and Recurrent Neural Networks
- Advanced Topics

Linear models + feature expansion recap



Finding good features (aka feature engineering) is a **highly non-trivial task**


Linear models + feature expansion recap



Finding good features (aka feature engineering) is a **highly non-trivial task**


Can we automate it?

Idea: add another model



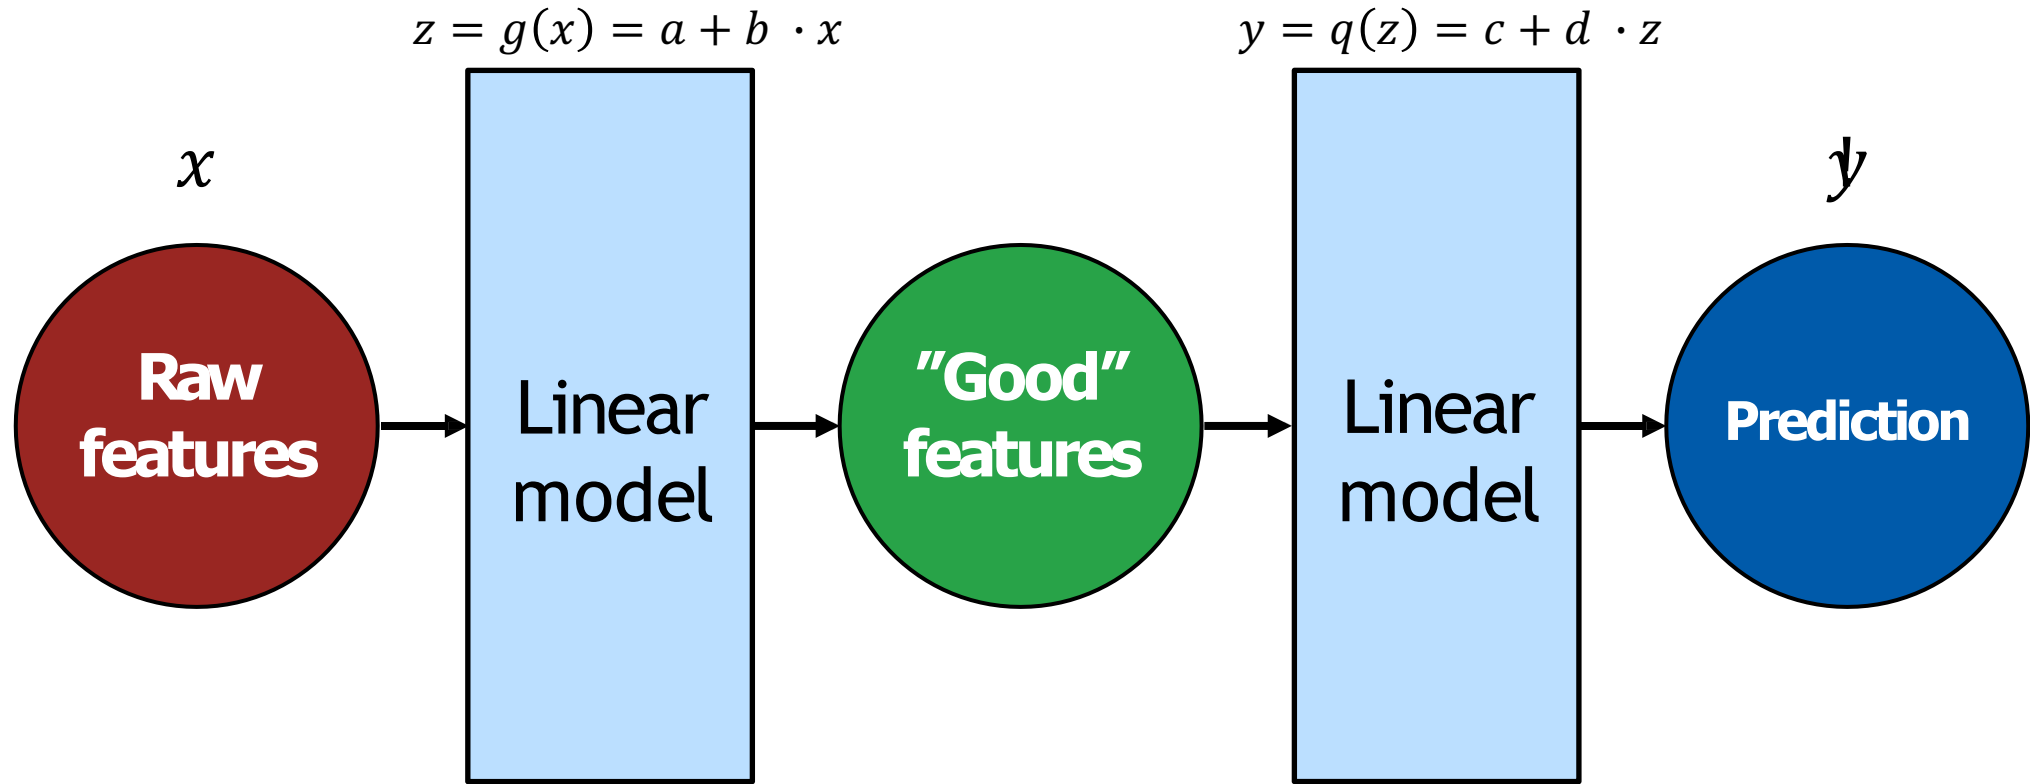
► Add another model

Idea: add another model



- ▶ Add another model
- ▶ Train everything simultaneously

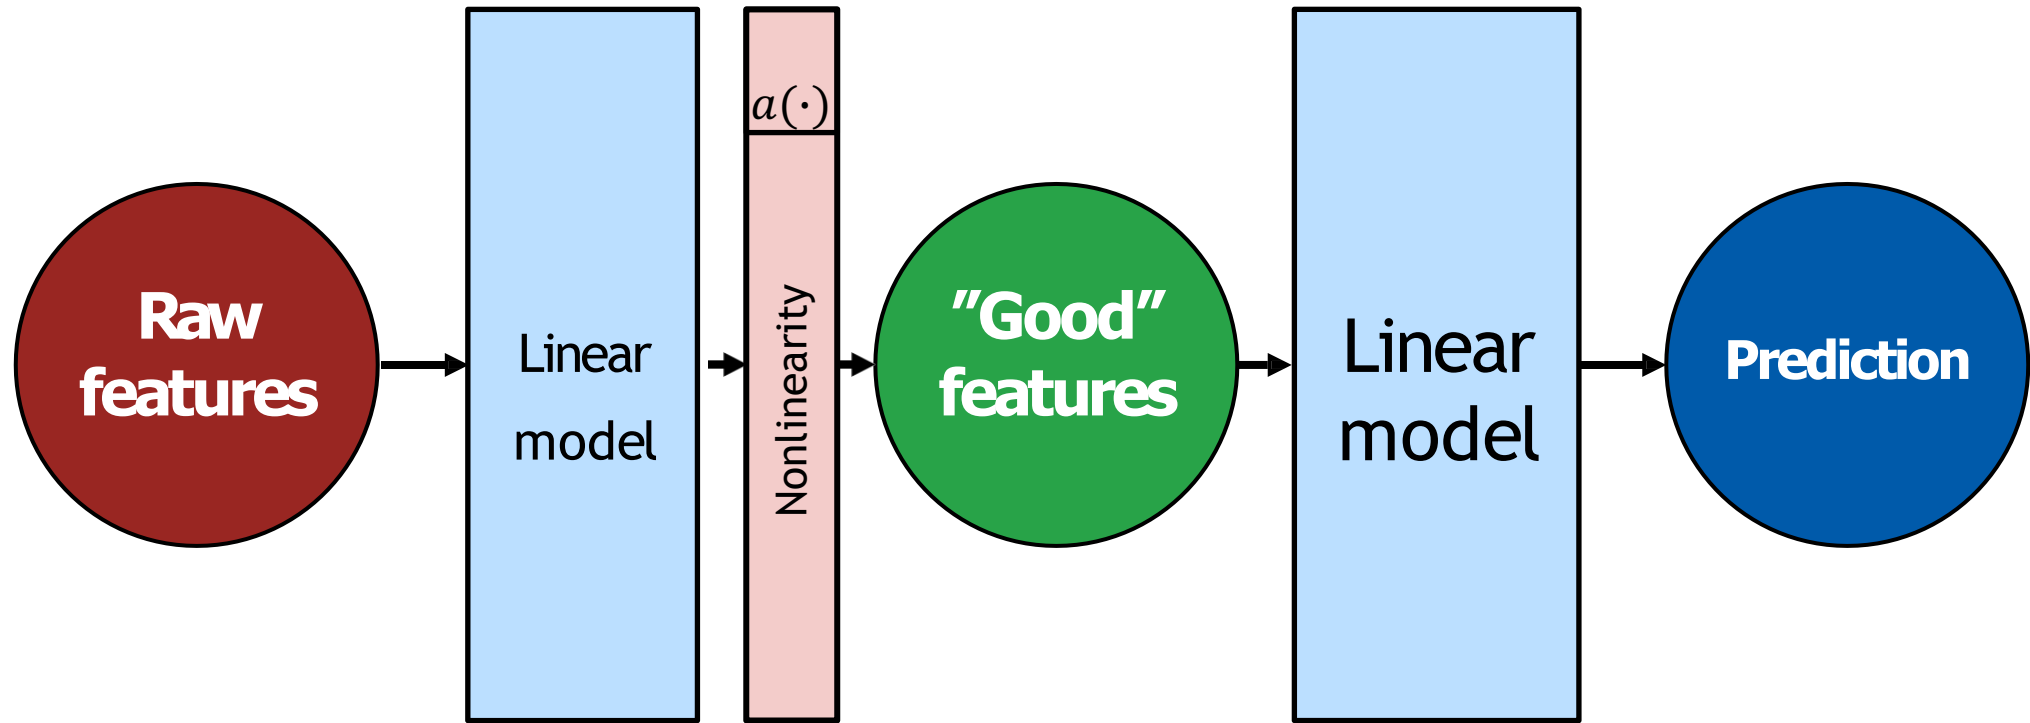
Can it be another linear model?



$$y = f(x) = q(g(x)) = c + d(a + b \cdot x) = c + d a + d b x$$

**NO! Two linear models results in a linear model.
We MUST add non linearities!!**


Fix: just introduce a nonlinearity




$$y = f(x) = q(a(g(x))) = c + d \cdot a(a + b \cdot x)$$

$a(\cdot)$ some **nonlinear** scalar function (applied elementwise)

Why they are called «neural»




A very common and widely-used model for the behaviour of neurons is based on the combination of a **linear function of several inputs**, then passed through a **non-linear** activation



Credits: http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture4.pdf


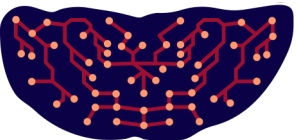
Are neural networks similar to the human brain?



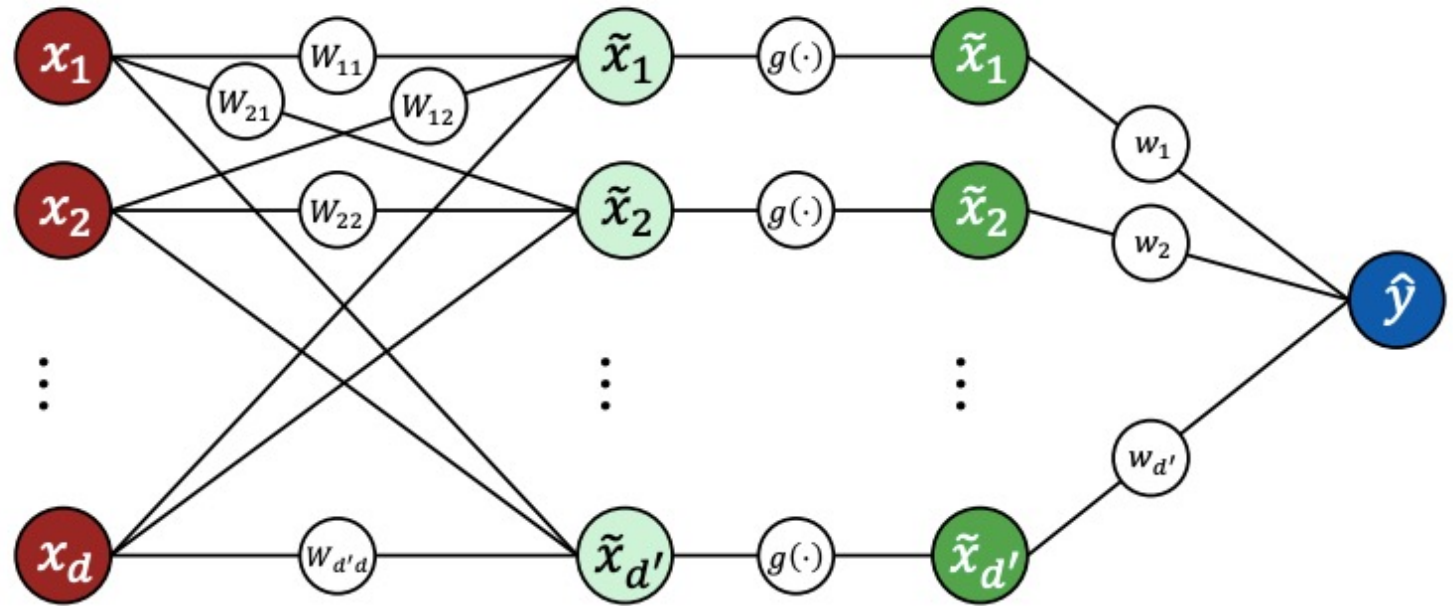
Similarities

1. As seen in the beginning, the **underlying model** of an artificial neuron was inspired the ones ones in biology to describe neuronal activations
2. Like in the brain, the combination and orchestration of **many single and relatively simple units** allow to build extremely powerful models

However, there are many important differences:

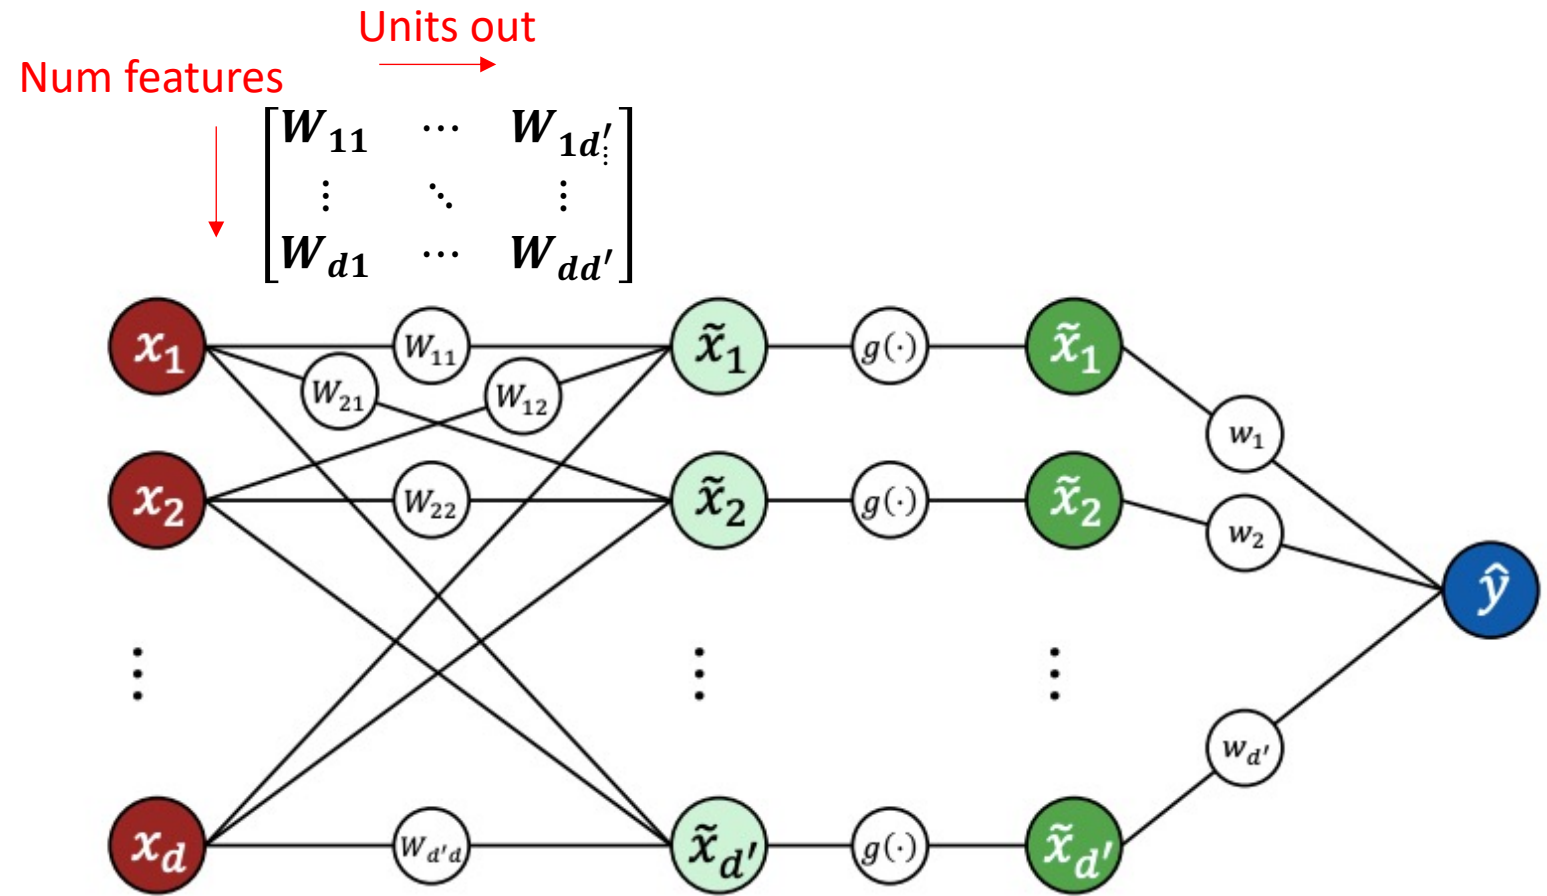
	Size	Topology	Activations	Efficiency	Learning
	~ 86 billions neurons	Very complex network with large-scale connections patterns	Based on biological physics, complex non linear dyn. system	Extremely low power consumptions and high efficiency	Able to generalize quickly and while learning topology changes
	~ 1-10k neurons, ~200 billions weights	Regular structure with links mostly between subsequent layers	Simple non-linear functions but of several types	Computations/learning requires large amount of energy	Backprop and optimization of non-convex function

Combining more inputs / neurons



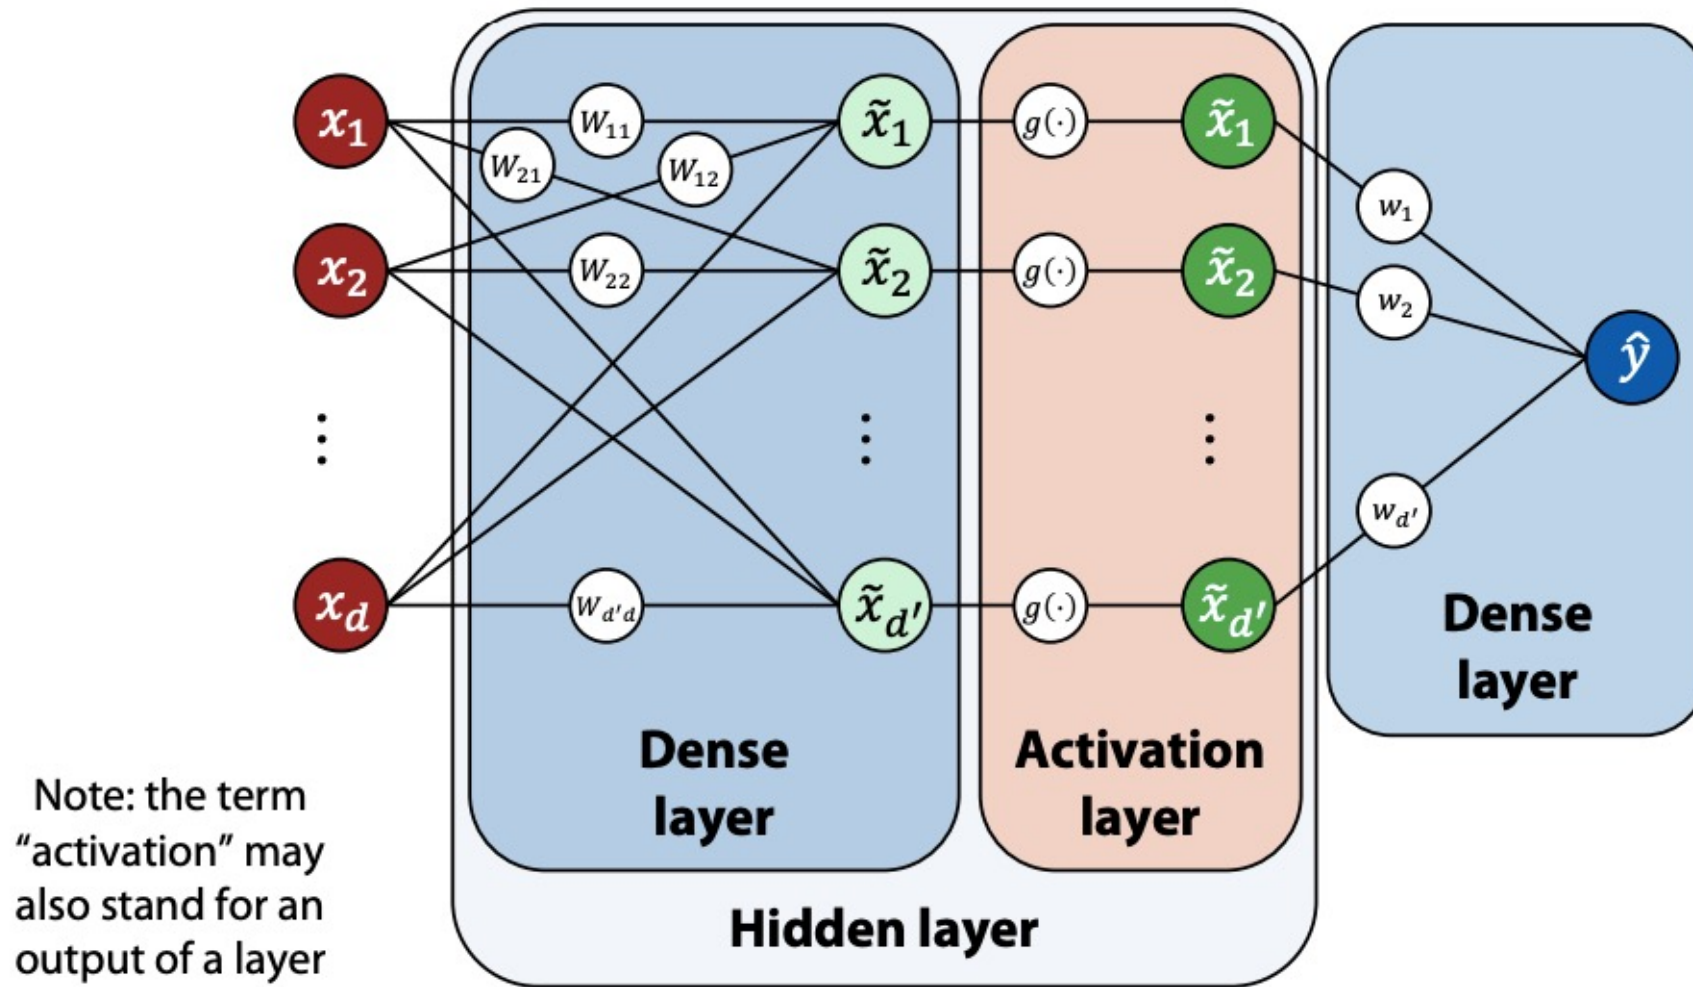
$$\hat{y} = w^T \tilde{x} = w^T g(Wx) = \sum_j \left[w_j g \left(\sum_i w_{ji} x_i \right) \right]$$

More details

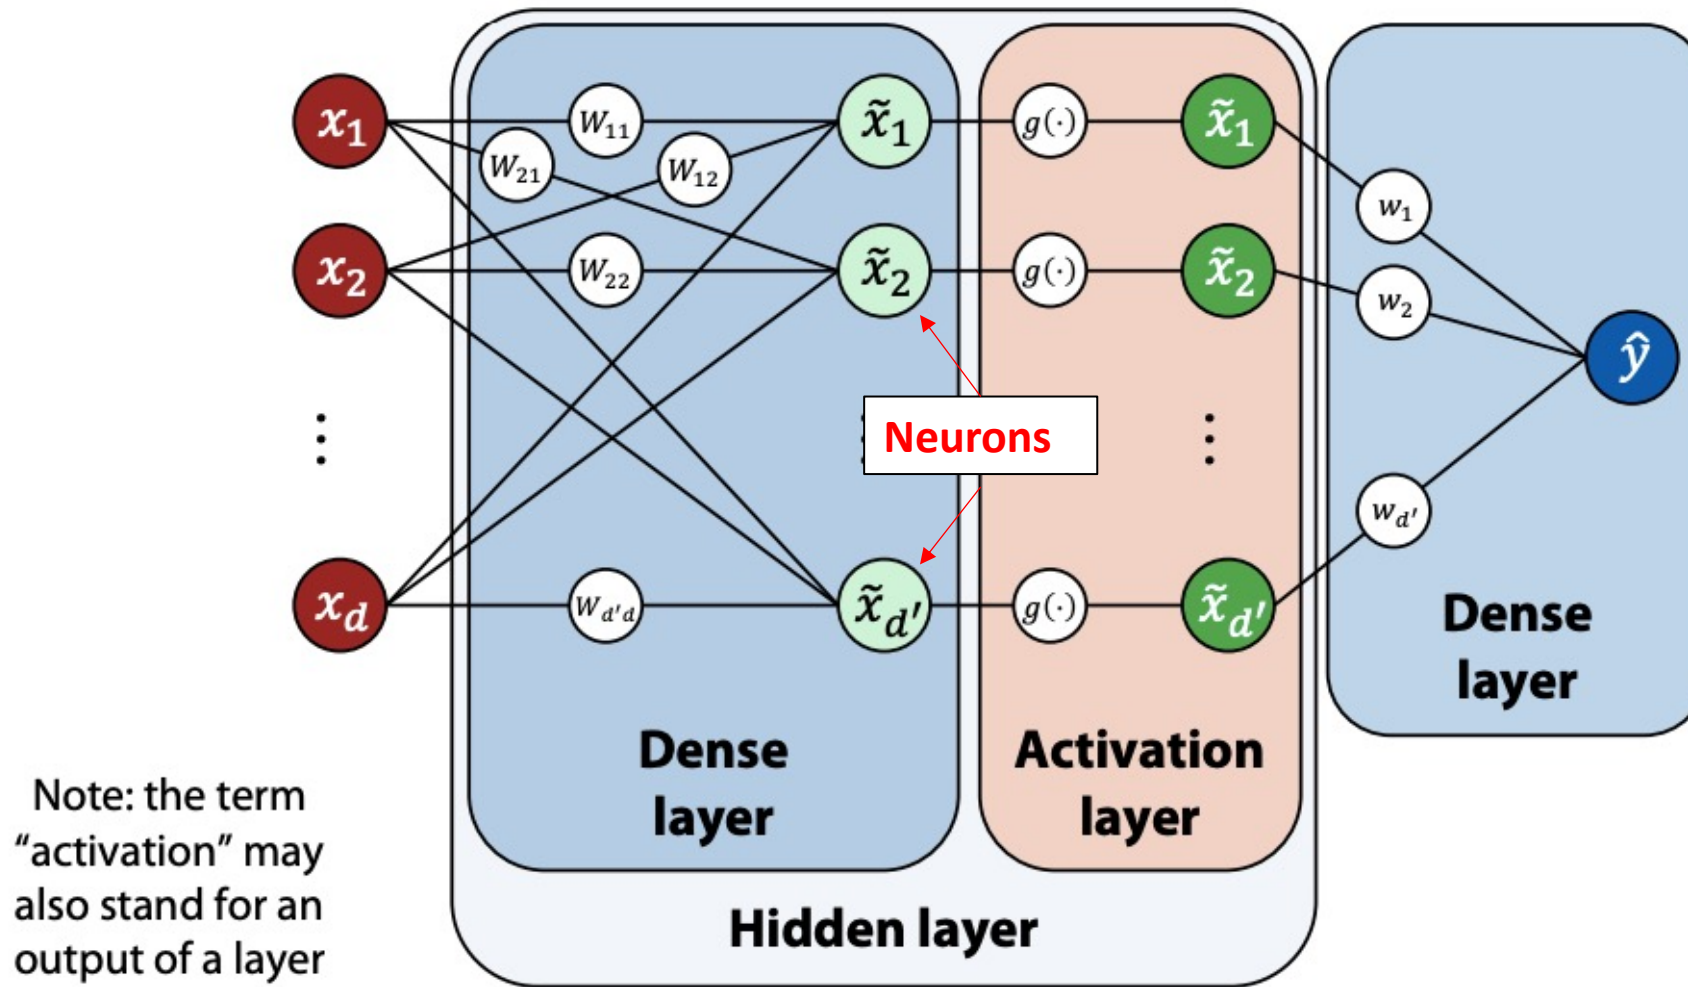


$$\hat{y} = w^T \tilde{x} = w^T g(Wx) = \sum_j \left[w_j g \left(\sum_i W_{ji} x_i \right) \right]$$


Terminology



Terminology: Simple MLP

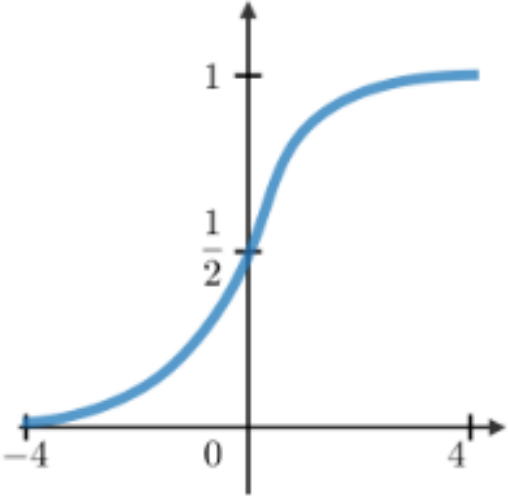
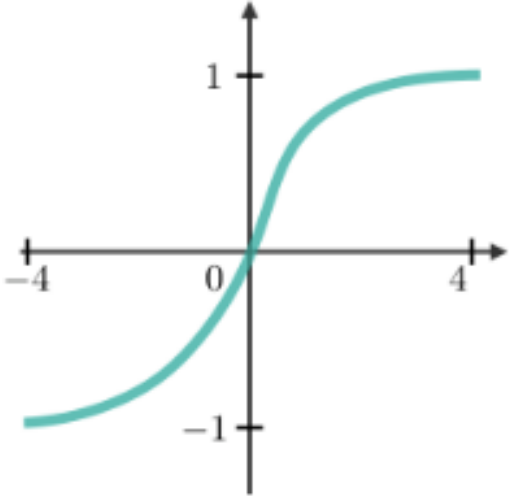
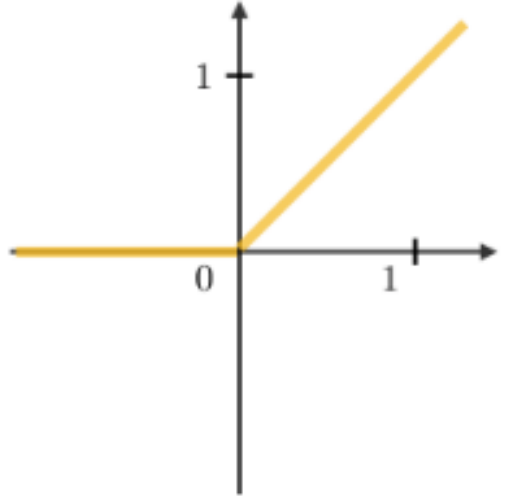


Terminology: Simple MLP



Credits: http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture4.pdf

Activation functions

Sigmoid	Tanh	RELU
$g(z) = \frac{1}{1 + e^{-z}}$	$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	$g(z) = \max(0, z)$
		

- Activation functions can be placed also before the output layer
- The output type determines which activation layer we need before the output

Universal Approximation Theorem (Cybenko, 1989)

“Neural Network can compute any function”

Universal Approximation Theorem

Universal approximation theorem for neural networks states that every continuous function that maps interval of real numbers can be approximated arbitrarily closely by a multi-layer perceptron with just one hidden layer.

Therefore

1. This does not mean we can compute any function but rather we can get approximations that is as good as we want
2. By increasing number of hidden units we can improve our approximation

BUT BE CAREFUL

We can get good approximations only for continuous functions: if functions are discontinuous (e.g. have sharp jumps) then approximation is difficult.

Universal Approximation Theorem (Cybenko, 1989)

“Neural Network can compute any function”

Universal Approximation Theorem

Universal approximation theorem for neural networks states that every continuous function that maps interval of real numbers can be approximated arbitrarily closely by a multi-layer perceptron with just one hidden layer.

Therefore

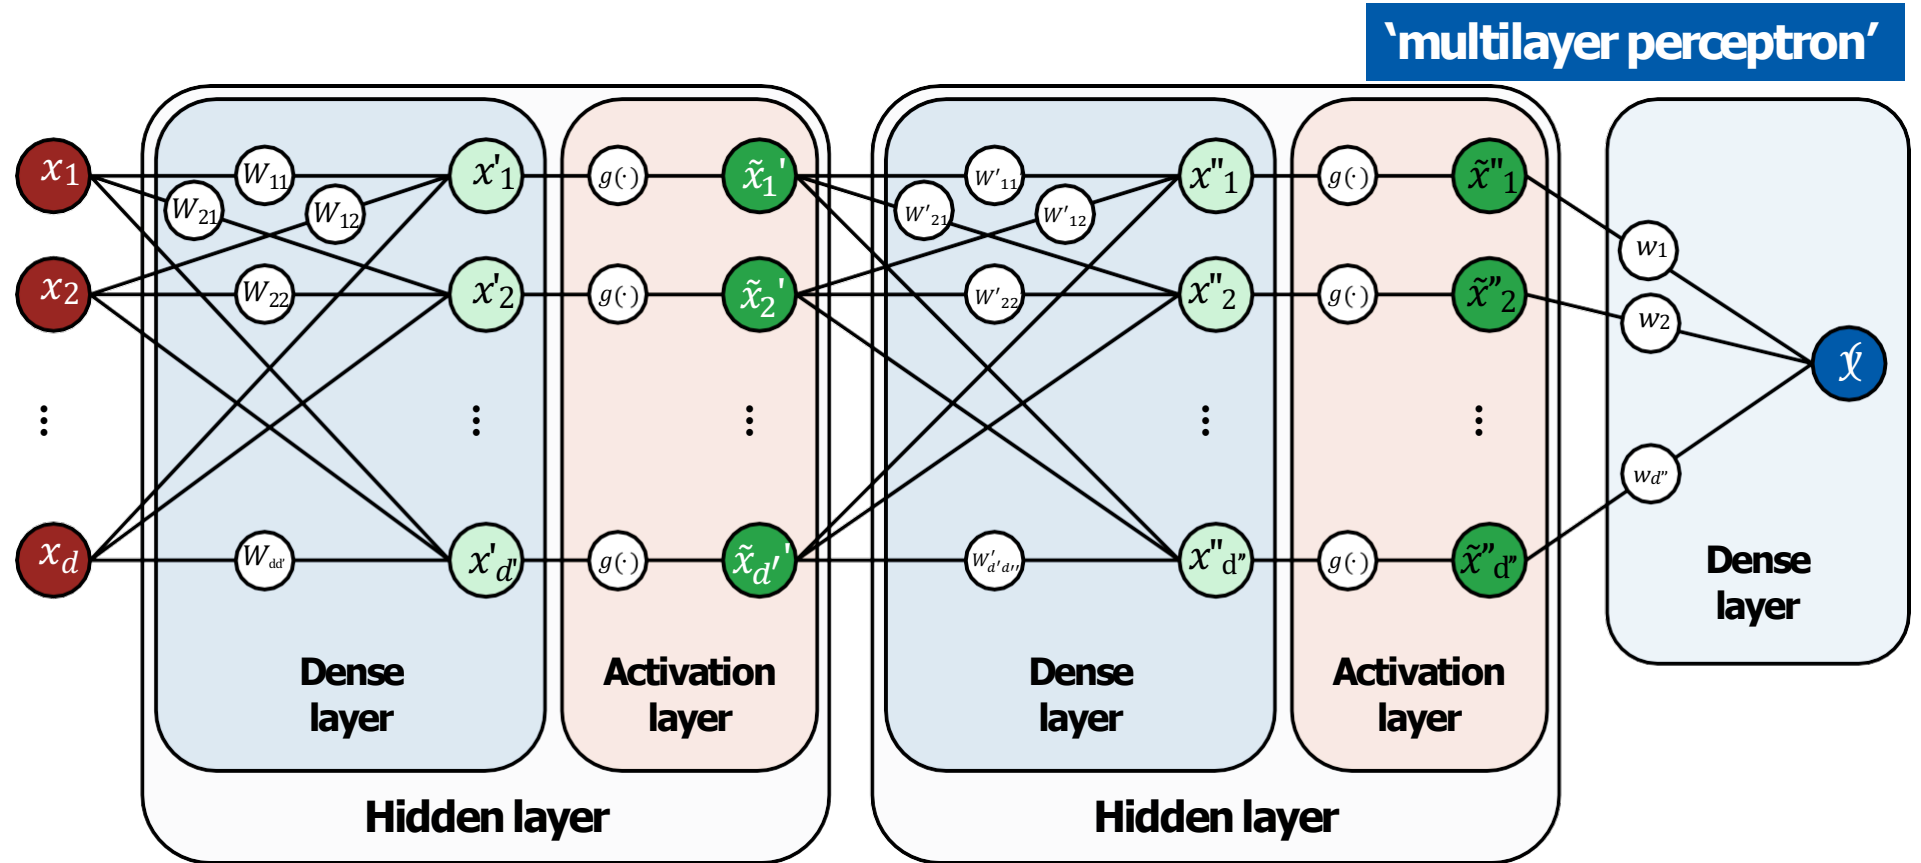
1. This does not mean we can compute any function but rather we can get approximations that is as good as we want
2. By increasing number of hidden units we can improve our approximation

BUT BE CAREFUL

We can get good approximations only for continuous functions: if functions are discontinuous (e.g. have sharp jumps) then approximation is difficult.


**Recently (Zhou et al, 2017 and Kidger and Lyons, 2020)
this has been extended for arbitrary depth!!!**

Deeper nets




In practice, stacking more hidden layers often reduces the number of neurons required to represent a given function

What are DEEP Neural Networks?



Why should we go deep?



1. Blue : Shallow model overfits after 20 million params
2. Red : Deep Model benefits from increasing number of layers
3. Increasing depth has more effect on learning of a model rather than increasing width of a layer

Goodfellow et al, Deep Learning, MIT Press, <http://www.deeplearningbook.org>

Why should we go deep?




Figure 6.6: Effect of depth. Empirical results showing that deeper networks generalize better when used to transcribe multidigit numbers from photographs of addresses. Data from Goodfellow *et al.* (2014d). The test set accuracy consistently increases with increasing depth. See figure 6.7 for a control experiment demonstrating that other increases to the model size do not yield the same effect.

Goodfellow et al, Deep Learning, MIT Press, <http://www.deeplearningbook.org>


Training a neural network

Training means find the value for the free parameters W and w^T . The process is very similar to what we have done for linear/polynomial regression, and it is formalized in terms of a **minimization problem**:

$$w = \arg \min L(w, \hat{y}_i, x_i, y_i)$$

where $L(w, \hat{y}_i, x_i, y_i)$ measures the distance of \hat{y} respect to its true value y .

$$\hat{y} = w^T \tilde{x} = w^T g(Wx) = \sum_j \left[w_j g \left(\sum_i W_{ji} x_i \right) \right]$$



Training a neural network

Training means find the value for the free parameters W and w^T . The process is very similar to what we have done for linear/polynomial regression, and it is formalized in terms of a **minimization problem**:

$$w = \arg \min L(w, \hat{y}_i, x_i, y_i)$$

where $L(w, \hat{y}_i, x_i, y_i)$ measures the distance of \hat{y} respect to its true value y .


$$\hat{y} = w^T \tilde{x} = \underbrace{w^T}_{\text{Parameters}} \underbrace{g(Wx)} = \sum_j \left[w_j \underbrace{g}_{\text{Activation Function}} \left(\sum_i W_{ji} x_i \right) \right]$$

Activation Function

The activation function generally makes the surface of the loss function **non-convex**

Fully Differentiable

The formulation is fully differentiable with respect to the weights: we can compute any derivative (see later backpropagation)



Training a neural network

- **Loss function $L(\mathbf{w})$**
 - Global error made by the model when compared with the target
 - Depends on the weights \mathbf{w} of the neural net

Training a neural network

- **Loss function $L(\mathbf{w})$**
 - Global error made by the model when compared with the target
 - Depends on the weights \mathbf{w} of the neural net


HOW TO FIND THE MINIMUM OF THIS FUNCTION?



Gradient descent


Gradient descent

- **Optimization algorithm:** numerical method to search the minimum (or maximum) of a function




Gradient descent

- **Optimization algorithm**: numerical method to search the minimum (or maximum) of a function
- **Iterative algorithm**: update the weights iteratively in the (opposite) direction of the gradient



Gradient descent

- **Optimization algorithm**: numerical method to search the minimum (or maximum) of a function
- **Iterative algorithm**: update the weights iteratively in the (opposite) direction of the gradient
- **Learning rate γ_n** : size of each step (can be different for each iteration)
 - Must be chosen wisely: if too large, it may skip over the minimum, if too small, convergence may be very slow




$$\mathbf{w}_{n+1} = \mathbf{w}_n - \gamma_n \nabla L(\mathbf{w}_n)$$

Gradient descent

- **Optimization algorithm**: numerical method to search the minimum (or maximum) of a function
- **Iterative algorithm**: update the weights iteratively in the (opposite) direction of the gradient
- **Learning rate γ_n** : size of each step (can be different for each iteration)
 - Must be chosen wisely: if too large, it may skip over the minimum, if too small, convergence may be very slow

Each training step (i.e. each time the weights are updated using all the training data) is called **EPOCH**



$$\mathbf{w}_{n+1} = \mathbf{w}_n - \gamma_n \nabla L(\mathbf{w}_n)$$

How to solve these equations?

Some packages can help us:



These libraries implement for us **backpropagation**, a technique which allows an intelligent sequential computation of the gradient.

The gradient is computed 'easily' exploiting the network structure of the algorithm → **computational graphs**